

LAPU-128

Instruction Set Reference

v0.6 (Draft) — November 2, 2025

Abstract

This document outlines the 128-bit instruction formats for LAPU-128, focused on complex arithmetic and vector descriptors. LAPU-128 is a small, focused ISA designed to perform complex tensor operations in an embedded environment. The following page shows the canonical XL, XC, XV, and XM encodings with 8-bit tick marks.

Core Instruction Formats (128-bit)

R-Type: Register-to-Register operations of either complex scalar or complex vector types

| R | | | | | | | | | | | | | | | | |
|---------------------|--------------------|-------------------|---------------|----------------|----------------|------------------|----|----|----|----|----|--------------------|-----|-----|-----|-----|
| 0 | 8 | 16 | 24 | 32 | 40 | 48 | 56 | 64 | 72 | 80 | 88 | 96 | 104 | 112 | 120 | 127 |
| opcode [127:120] | subop [119:112] | flags [111:96] | rd [95:93] | rs1 [92:90] | rs2 [89:87] | imm16 [86:71] | | | | | | reserved [70:0] | | | | |

I-type: Immediate operations of just complex scalars

| I | | | | | | | | | | | | | | | | |
|---------------------|--------------------|-------------------|---------------|----------------|----|----|----|----|----|----|----|------------------|-----|-----|-----|-----|
| 0 | 8 | 16 | 24 | 32 | 40 | 48 | 56 | 64 | 72 | 80 | 88 | 96 | 104 | 112 | 120 | 127 |
| opcode [127:120] | subop [119:112] | flags [111:96] | rd [95:93] | rs1 [92:90] | | | | | | | | imm_90 [89:0] | | | | |

J (conditional jump, 128-bit descriptor)

| JX | | | | | | | | | | | | | | | | |
|------------------|--------------------|---------------------|----------------|----|----------------------------|----|----|----|----|----|----|--------------------|-----|-----|-----|-----|
| 0 | 8 | 16 | 24 | 32 | 40 | 48 | 56 | 64 | 72 | 80 | 88 | 96 | 104 | 112 | 120 | 127 |
| opc [127:120] | subop [119:112] | flags16 [111:96] | rs1 [95:93] | | offs33 (PC_rel) [92:60] | | | | | | | reserved [59:0] | | | | |

S-type: matrix-bank load/store (scalar & vector), 128-bit

| S-type (stride) | | | | | | | | | | | | | | | | |
|------------------|--------------------|---------------------|---------------|-----------------|----------------|----------------|----|----|----|----|----|---------------|-----|-----|-----|-----|
| 0 | 8 | 16 | 24 | 32 | 40 | 48 | 56 | 64 | 72 | 80 | 88 | 96 | 104 | 112 | 120 | 127 |
| opc [127:120] | subop [119:112] | FLAGS16 [111:96] | rd [95:93] | mbid [92:89] | i16 [88:73] | j16 [72:57] | | | | | | rsv [56:0] | | | | |

Register Layout

Architectural Registers (Summary)

| Class | Names | Width / Elements | Notes |
|------------------|----------|---|--|
| Scalar (complex) | $s0..s7$ | 128 b each (complex Q32.32 + Q32.32) | $s0$ is hard-wired to 0. $s1$ is the conventional branch predicate (0/1). |
| Vector (complex) | $v0..v7$ | VLEN elements; each element 128 b complex | $v0$ is hard-wired to all-zeros. Vector ops always operate on all VLEN elements. |

Vector Length

VLEN is a hardware/HDL parameter fixed at synthesis time. It is constant at runtime. All vector instructions operate over the entire range [0, VLEN – 1].

Complex Number Format (Q32.32 + Q32.32)

Each scalar register and each vector element encodes a complex value (Re, Im) in fixed point:

$$\text{Re, Im} \in \text{Q32.32 two's complement} \Rightarrow x_{\text{real}} = \frac{X_{\text{int}}}{2^{32}}, x_{\text{imag}} = \frac{Y_{\text{int}}}{2^{32}}.$$

The 128-bit complex is stored little-endian in memory with **Re at the lower address** and **Im at the higher address**. Each half (Re or Im) is a 64-bit two's-complement fixed-point integer with 32 integer bits and 32 fractional bits.

Endianness

Instruction words (128 b) and data are little-endian. For complex numbers in memory: bytes for Re precede bytes for Im.

Zero Registers

The following are architecturally fixed to zero and never written:

$$s0 \equiv 0 \quad (\text{complex zero}), \quad v0[i] \equiv 0 \quad \forall i \in [0, \text{VLEN} - 1].$$

Matrix Banks

There will be 4 matrix banks to choose from. Each matrix bank is a square matrix of side N times VLEN where N is some predefined integer greater than 1. The matrix bank will be accessed either through scalar (individual) read/writes or vector read/writes. Vector read/writes will access vectors in either row or major column order and will have no overlap between each other in either mode.

Instruction Semantics

R-type — Register-to-Register (complex)

OPCODE: 0x01

Description

These are register to register operations involve either two vectors, two scalars, or a vector and a scalar. To determine mapping check bit position [97:96] under flags. Additionally each subop code range will be defined per mapping

$$\begin{aligned} f(s_1, s_2) &\mapsto s \in S & 00 \\ f(\mathbf{v}_1, \mathbf{v}_2) &\mapsto \mathbf{v} \in V & 01 \\ f(\mathbf{v}_1, \mathbf{v}_2) &\mapsto s \in S & 10 \\ f(\mathbf{v}, s) &\mapsto \mathbf{v} \in V & 11 \end{aligned}$$

Encoding notes: The mapping selector lives in `flags[97:96]`. Field `imm16` is currently unused and must be zero. All reserved bits [70:0] must be zero.

Scalar ops (unary and binary)

Table 2: Scalar register ops (`s*`): $S \rightarrow S$ and $S \times S \rightarrow S$

| Mnemonic | subop | Operands | Effect | Notes |
|--|-------|----------|--|--|
| <i>Unary: $S \rightarrow S$</i> | | | | |
| cneg | 0x00 | d, a | $d \leftarrow -a$ | Two's-complement both halves. |
| conj | 0x01 | d, a | $d \leftarrow \text{conj}(a)$ | Negate imaginary half. |
| csqrt | 0x02 | d, a | $d \leftarrow \sqrt{a}$ | Principal root; widen, then truncate to Q32.32+Q32.32. |
| cabs2 | 0x03 | d, a | $d_{\text{re}} \leftarrow \Re(a)^2 + \Im(a)^2$, $d_{\text{im}} \leftarrow 0$ | Magnitude ² ; widen then truncate. |
| cabs | 0x04 | d, a | $d_{\text{re}} \leftarrow \sqrt{\Re(a)^2 + \Im(a)^2}$, $d_{\text{im}} \leftarrow 0$ | Fixed-point $\sqrt{\cdot}$; truncating. |
| creal | 0x05 | d, a | $d_{\text{re}} \leftarrow \Re(a)$, $d_{\text{im}} \leftarrow 0$ | Extract real. |
| cimag | 0x06 | d, a | $d_{\text{re}} \leftarrow \Im(a)$, $d_{\text{im}} \leftarrow 0$ | Extract imaginary to real half. |
| crecip | 0x07 | d, a | $d \leftarrow 1 \div a$ | $(\bar{a})/ a ^2$; if $a=0$ then $d:=0$. |
| <i>Binary: $S \times S \rightarrow S$</i> | | | | |
| cadd | 0x08 | d, a, b | $d \leftarrow a + b$ | Truncating Q32.32+Q32.32. |
| csub | 0x09 | d, a, b | $d \leftarrow a - b$ | Truncating. |
| cmul | 0x0A | d, a, b | $d \leftarrow a \times b$ | Widen internally, truncate to Q32.32+Q32.32. |
| cdiv | 0x0B | d, a, b | $d \leftarrow a \div b$ | $(a \bar{b})/ b ^2$; if $ b =0$ then $d:=0$. |

Vector → Vector

Table 3: Lane-wise vector ops (`v*`): $V \rightarrow V$ and $V \times V \rightarrow V$

| Mnemonic | subop | Operands | Effect | Notes |
|----------|-------|------------|---------------------------------------|-----------------------------|
| vadd | 0x00 | vD, vA, vB | $vD[i] \leftarrow vA[i] + vB[i]$ | Saturating per lane. |
| vsub | 0x01 | vD, vA, vB | $vD[i] \leftarrow vA[i] - vB[i]$ | Saturating per lane. |
| vmul | 0x02 | vD, vA, vB | $vD[i] \leftarrow vA[i] \times vB[i]$ | Complex lane-wise multiply. |

| Mnemonic | subop | Operands | Effect | Notes |
|----------|-------|------------|---------------------------------------|---|
| vdiv | 0x04 | vD, vA, vB | $vD[i] \leftarrow vA[i] \div vB[i]$ | $(a\bar{b})/ b ^2$; if $ b =0$ then lane:=0. |
| vconj | 0x05 | vD, vA | $vD[i] \leftarrow \text{conj}(vA[i])$ | Lane-wise conjugate. |

Vector / Vector → Scalar (reductions)

Table 4: Reductions to scalar: $V \rightarrow S$ and $V \times V \rightarrow S$

| Mnemonic | subop | Operands | Effect | Notes |
|----------|-------|------------------------|--|---------------------------------|
| dotu | 0x01 | sD, \bar{A}, \bar{B} | $sD \leftarrow \sum_{i=0}^{\text{VLEN}-1} \bar{A}[i] \bar{B}[i]$ | Complex dot (no conjugation). |
| iamax | 0x02 | sD, vA | $sD \leftarrow \arg \max_i vA[i] $ | Index in sD real half; imag:=0. |
| sum | 0x03 | sD, \bar{A} | $sD \leftarrow \sum_{i=0}^{\text{VLEN}-1} \bar{A}[i]$ | Complex sum; reduces to scalar. |

Vector × Scalar → Vector (broadcast per lane)

Table 5: Vector–scalar broadcast ops: $V \times S \rightarrow V$

| Mnemonic | subop | Operands | Effect | Notes |
|----------|-------|------------|------------------------------------|--|
| vsadd | 0x18 | vD, vA, sB | $vD[i] \leftarrow vA[i] + sB$ | Broadcast add; saturating per lane. |
| vssub | 0x19 | vD, vA, sB | $vD[i] \leftarrow vA[i] - sB$ | Broadcast sub (vector minus scalar); saturating per lane. |
| vsmul | 0x1A | vD, vA, sB | $vD[i] \leftarrow vA[i] \times sB$ | Complex lane-wise multiply by complex scalar; widen then truncate. |
| vsdiv | 0x1B | vD, vA, sB | $vD[i] \leftarrow vA[i] \div sB$ | $(a\bar{b})/ b ^2$ per lane; if $ sB =0$ then lane:=0. |

I-type — Immediate (scalars only)

OPCODE: 0x02

Description

`imm_90` is a complex number split into two 45-bit **signed** fixed-point halves (Q22.23, two's complement), packed as:

$$\text{Re} \rightarrow \text{imm_90}[89:45], \quad \text{Im} \rightarrow \text{imm_90}[44:0].$$

Table 6: I-type: Immediate operations (scalar complex)

| Mnemonic | subop | Operands | Effect | Notes |
|---------------------|-------|--------------|--------------------------------|--|
| <code>cloadi</code> | 0x00 | sD, cIMM | $sD \leftarrow cIMM$ | cIMM packed in <code>imm_90</code> (Re/Im per above). |
| <code>cadd_i</code> | 0x01 | sD, sA, cIMM | $sD \leftarrow sA + cIMM$ | Saturating per scalar; Q32.32 truncation as needed. |
| <code>cmul_i</code> | 0x02 | sD, sA, cIMM | $sD \leftarrow sA \times cIMM$ | Widen internally, clamp/truncate to Q32.32. |
| <code>csub_i</code> | 0x03 | sD, sA, cIMM | $sD \leftarrow sA - cIMM$ | Saturating; truncation semantics match <code>csub.c</code> . |
| <code>cdiv_i</code> | 0x04 | sD, sA, cIMM | $sD \leftarrow sA \div cIMM$ | $(sA \overline{cIMM}) / cIMM ^2$; if $ cIMM =0$ then $sD:=0$. |

J-type — Conditional Jump

OPCODE: 0x03

Table 7: J-type: Conditional jump (single predicate)

| Mnemonic | subop | Operands | Effect | Notes |
|-------------------|-------|----------|---|--|
| <code>jrel</code> | 0x00 | offs33 | If $s1 \neq 0$: $\text{PC} \leftarrow \text{PC} + \text{offs33}$ (instruction-relative). | Encoding: <code>offs33</code> is <i>signed</i> two's complement in <i>instruction units</i> (128-bit words); base is the address of <i>this</i> instruction. Field <code>rs1</code> must be 001b (predicated on <code>s1</code>). All <code>flags16</code> and reserved bits must be zero. |

S-type — Matrix-bank Vector Load/Store (stride implicit)

OPCODE: 0x04

Table 8: S-type: Matrix-bank vector load/store

| Mnemonic | subop | Operands | Effect | Notes |
|----------|-------|---------------------|---|--|
| vld | 0x00 | vD, mbid, rc, idx16 | Load into vD the sequence: if $rc=0$: $(r=idx16, c=0..L-1)$, if $rc=1$: $(r=0..L-1, c=idx16)$, where $L = \text{len16}$ if nonzero, else $L = \text{VLEN}$. | Encoding: FLAGS16.rc is bit 111 (MSB of FLAGS16); rc =0 means row, rc =1 means column. Fields idx16 and idy16 are unsigned 16-bit ; Write as vld.rm or vld.cm to select row or column major order. Elements are 128-bit complex (Re then Im). |
| vst | 0x01 | vS, mbid, rc, idx16 | Store from vS to the same address pattern as vld. | Same encoding and field conventions as vld. |
| sld.xy | 0x02 | sD, mbid, x16, y16 | Load the single element at coordinates $(r=y16, c=x16)$ from matrix-bank mbid into scalar register sD. | Coordinates are 0-based unsigned 16-bit . Mapping: $x16 \rightarrow i16[88:73]$, $y16 \rightarrow j16[72:57]$. Element size is one 128-bit complex (Re then Im). Out-of-bounds coordinates trap. |
| sst.xy | 0x03 | sS, mbid, x16, y16 | Store scalar sS to the element at $(r=y16, c=x16)$ in matrix-bank mbid. | Same addressing and trapping rules as sld.xy. Mapping: $x16 \rightarrow i16[88:73]$, $y16 \rightarrow j16[72:57]$. |

Global encoding rules. Unless otherwise specified, all **reserved** fields are **must-be-zero** and all currently **unused** fields are encoded as **zero**.