

LAPU-128

Instruction Set Reference

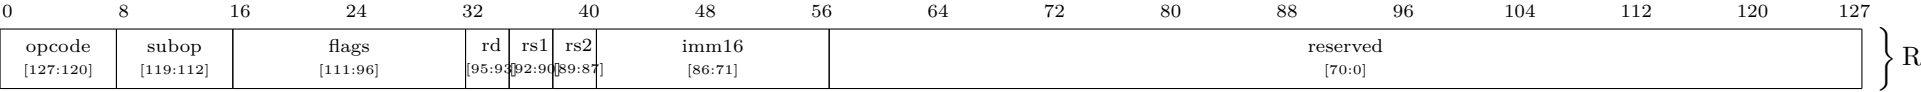
v0.6 (Draft) — October 14, 2025

Abstract

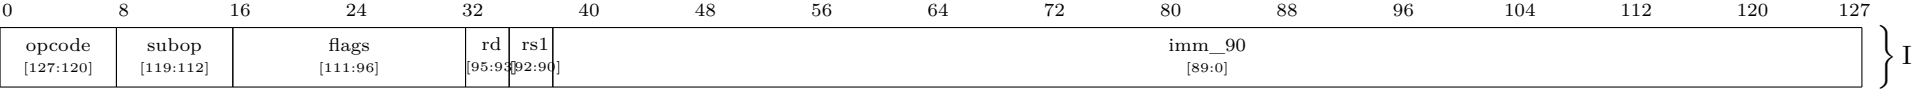
This document outlines the 128-bit instruction formats for LAPU-128, focused on complex arithmetic and vector descriptors. LAPU-128 is a small, focused ISA designed to perform complex tensor operations in an embedded environment. The following page shows the canonical XL, XC, XV, and XM encodings with 8-bit tick marks.

Core Instruction Formats (128-bit)

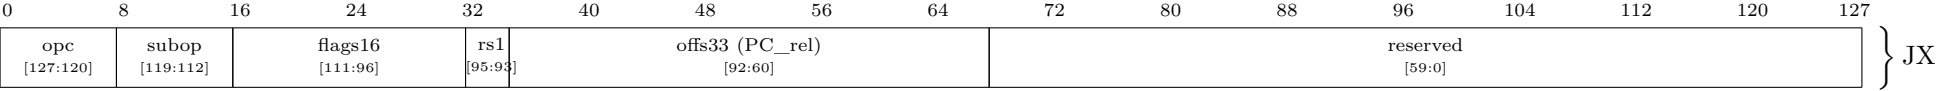
R-Type: Register-to-Register operations of either complex scalar or complex vector types



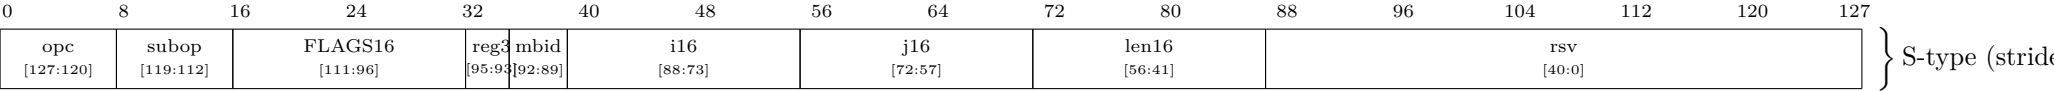
I-type: Immediate operations of just complex scalars



J (conditional jump, 128-bit descriptor)



S-type: matrix-bank load/store (scalar & vector), 128-bit



Register Layout

Architectural Registers (Summary)

Class	Names	Width / Elements	Notes
Scalar (complex)	$s0..s7$	128 b each (complex Q32.32 + Q32.32)	$s0$ is hard-wired to 0 . $s1$ is the conventional branch predicate (0/1).
Vector (complex)	$v0..v7$	VLEN elements; each element 128 b complex	$v0$ is hard-wired to all-zeros . Vector ops always operate on all VLEN elements .

Vector Length

VLEN is a hardware/HDL parameter fixed at synthesis time. It is constant at runtime. All vector instructions operate over the entire range $[0, \text{VLEN} - 1]$.

Complex Number Format (Q32.32 + Q32.32)

Each scalar register and each vector element encodes a complex value (Re, Im) in fixed point:

$$\text{Re, Im} \in \text{Q32.32 two's complement} \Rightarrow x_{\text{real}} = \frac{X_{\text{int}}}{2^{32}}, \quad x_{\text{imag}} = \frac{Y_{\text{int}}}{2^{32}}.$$

The 128-bit complex is stored little-endian in memory with **Re at the lower address** and **Im at the higher address**. Each half (Re or Im) is a 64-bit two's-complement fixed-point integer with 32 integer bits and 32 fractional bits.

Endianness

Instruction words (128 b) and data are little-endian. For complex numbers in memory: bytes for Re precede bytes for Im.

Zero Registers

The following are architecturally fixed to zero and never written:

$$s0 \equiv 0 \quad (\text{complex zero}), \quad v0[i] \equiv 0 \quad \forall i \in [0, \text{VLEN} - 1].$$

Matrix Banks

There will be 4 matrix banks to choose from. Each matrix bank is a square matrix of side N times VLEN where N is some predefined integer greater than 1. The matrix bank will be accessed either through scalar (individual) read/writes or vector read/writes. Vector read/writes will access vectors in either row or major column order and will have no overlap between each other in either mode.

Instruction Semantics

R-type — Register-to-Register (complex)

OPCODE: 0x01

Description

These are register to register operations involve either two vectors, two scalars, or a vector and a scalar. To determine mapping check bit position [97:96] under flags. Additionally each subop code range will be defined per mapping

$$\begin{aligned} f(s_1, s_2) &\mapsto s \in S & 00 \\ f(\mathbf{v}_1, \mathbf{v}_2) &\mapsto \mathbf{v} \in V & 01 \\ f(\mathbf{v}_1, \mathbf{v}_2) &\mapsto s \in S & 10 \\ f(\mathbf{v}, s) &\mapsto \mathbf{v} \in V & 11 \end{aligned}$$

Encoding notes: The mapping selector lives in `flags[97:96]`. Field `imm16` is currently unused and **must be zero**. All **reserved** bits [70:0] **must be zero**.

Scalar ops (unary and binary)

Table 2: Scalar register ops (**s***): $S \rightarrow S$ and $S \times S \rightarrow S$

Mnemonic	subop	Operands	Effect	Notes
<i>Unary: $S \rightarrow S$</i>				
<code>cneg</code>	0x00	d, a	$d \leftarrow -a$	Two's-complement both halves.
<code>conj</code>	0x01	d, a	$d \leftarrow \text{conj}(a)$	Negate imaginary half.
<code>csqrt</code>	0x02	d, a	$d \leftarrow \sqrt{a}$	Principal root; widen, then truncate to Q32.32+Q32.32.
<code>cabs2</code>	0x03	d, a	$d_{\text{re}} \leftarrow \Re(a)^2 + \Im(a)^2, d_{\text{im}} \leftarrow 0$	Magnitude ² ; widen then truncate.
<code>cabs</code>	0x04	d, a	$d_{\text{re}} \leftarrow \sqrt{\Re(a)^2 + \Im(a)^2}, d_{\text{im}} \leftarrow 0$	Fixed-point $\sqrt{\cdot}$; truncating.
<code>creal</code>	0x05	d, a	$d_{\text{re}} \leftarrow \Re(a), d_{\text{im}} \leftarrow 0$	Extract real.
<code>cimag</code>	0x06	d, a	$d_{\text{re}} \leftarrow \Im(a), d_{\text{im}} \leftarrow 0$	Extract imaginary to real half.
<code>crecip</code>	0x07	d, a	$d \leftarrow 1 \div a$	$(\bar{a})/ a ^2$; if $a=0$ then $d:=0$.
<i>Binary: $S \times S \rightarrow S$</i>				
<code>cadd</code>	0x08	d, a, b	$d \leftarrow a + b$	Truncating Q32.32+Q32.32.
<code>csub</code>	0x09	d, a, b	$d \leftarrow a - b$	Truncating.
<code>cmul</code>	0x0A	d, a, b	$d \leftarrow a \times b$	Widen internally, truncate to Q32.32+Q32.32.
<code>cdiv</code>	0x0B	d, a, b	$d \leftarrow a \div b$	$(a\bar{b})/ b ^2$; if $ b =0$ then $d:=0$.
<code>cmaxabs</code>	0x0C	d, a, b	$d \leftarrow \arg \max_{x \in \{a,b\}} x $	Ties pick a .
<code>cminabs</code>	0x0D	d, a, b	$d \leftarrow \arg \min_{x \in \{a,b\}} x $	Ties pick a .
<code>cmplt.re</code>	0x0E	d, a, b	$d_{\text{re}} \leftarrow 1 \text{ if } \Re(a) < \Re(b) \text{ else } 0; d_{\text{im}} \leftarrow 0$	Ignores imaginary parts of a and b .
<code>cmpgt.re</code>	0x0F	d, a, b	$d_{\text{re}} \leftarrow 1 \text{ if } \Re(a) > \Re(b) \text{ else } 0; d_{\text{im}} \leftarrow 0$	Ignores imaginary parts of a and b .
<code>cmple.re</code>	0x10	d, a, b	$d_{\text{re}} \leftarrow 1 \text{ if } \Re(a) \leq \Re(b) \text{ else } 0; d_{\text{im}} \leftarrow 0$	Ignores imaginary parts of a and b .

Vector → Vector

Table 3: Lane-wise vector ops ($\mathbf{v}*$): $\mathbf{V} \rightarrow \mathbf{V}$ and $\mathbf{V} \times \mathbf{V} \rightarrow \mathbf{V}$

Mnemonic	subop	Operands	Effect	Notes
vadd	0x00	vD, vA, vB	$\mathbf{vD}[i] \leftarrow \mathbf{vA}[i] + \mathbf{vB}[i]$	Saturating per lane.
vsub	0x01	vD, vA, vB	$\mathbf{vD}[i] \leftarrow \mathbf{vA}[i] - \mathbf{vB}[i]$	Saturating per lane.
vmul	0x02	vD, vA, vB	$\mathbf{vD}[i] \leftarrow \mathbf{vA}[i] \times \mathbf{vB}[i]$	Complex lane-wise multiply.
vmac	0x03	vD, vA, vB	$\mathbf{vD}[i] \leftarrow \mathbf{vD}[i] + \mathbf{vA}[i] \times \mathbf{vB}[i]$	Fused complex MAC per lane.
vdiv	0x04	vD, vA, vB	$\mathbf{vD}[i] \leftarrow \mathbf{vA}[i] \div \mathbf{vB}[i]$	$(a\bar{b})/ b ^2$; if $ b =0$ then lane:=0.
vconj	0x05	vD, vA	$\mathbf{vD}[i] \leftarrow \text{conj}(\mathbf{vA}[i])$	Lane-wise conjugate.

Vector / Vector → Scalar (reductions)

Table 4: Reductions to scalar: $\mathbf{V} \rightarrow \mathbf{S}$ and $\mathbf{V} \times \mathbf{V} \rightarrow \mathbf{S}$

Mnemonic	subop	Operands	Effect	Notes
dotc	0x00	sD, vA, vB	$sD \leftarrow \sum_{i=0}^{\text{VLEN}-1} \text{conj}(vA[i]) \cdot vB[i]$	Reduce to scalar sD.
dotu	0x01	sD, \bar{A} , \bar{B}	$sD \leftarrow \sum_{i=0}^{\text{VLEN}-1} \bar{A}[i] \bar{B}[i]$	Complex dot (no conjugation).
iamax	0x02	sD, vA	$sD \leftarrow \arg \max_i vA[i] $	Index in sD real half; imag:=0.
sum	0x03	sD, \bar{A}	$sD \leftarrow \sum_{i=0}^{\text{VLEN}-1} \bar{A}[i]$	Complex sum; reduces to scalar.
asum	0x04	sD, \bar{A}	$sD_{\text{re}} \leftarrow \sum_{i=0}^{\text{VLEN}-1} \bar{A}[i] , sD_{\text{im}} \leftarrow 0$	Sum of magnitudes (real result).

Vector × Scalar → Vector (broadcast per lane)

Table 5: Vector-scalar broadcast ops: $\mathbf{V} \times \mathbf{S} \rightarrow \mathbf{V}$

Mnemonic	subop	Operands	Effect	Notes
vsadd	0x18	vD, vA, sB	$\mathbf{vD}[i] \leftarrow \mathbf{vA}[i] + sB$	Broadcast add; saturating per lane.
vssub	0x19	vD, vA, sB	$\mathbf{vD}[i] \leftarrow \mathbf{vA}[i] - sB$	Broadcast sub (vector minus scalar); saturating per lane.
vsmul	0x1A	vD, vA, sB	$\mathbf{vD}[i] \leftarrow \mathbf{vA}[i] \times sB$	Complex lane-wise multiply by complex scalar; widen then truncate.
vsdiv	0x1B	vD, vA, sB	$\mathbf{vD}[i] \leftarrow \mathbf{vA}[i] \div sB$	$(a\bar{b})/ b ^2$ per lane; if $ sB =0$ then lane:=0.

I-type — Immediate (scalars only)

OPCODE: 0x02

Description

imm_90 is a complex number split into two 45-bit **signed** fixed-point halves (Q22.23, two's complement), packed as:

$$\text{Re} \rightarrow \text{imm_90}[44:0], \quad \text{Im} \rightarrow \text{imm_90}[89:45].$$

Table 6: I-type: Immediate operations (scalar complex)

Mnemonic	subop	Operands	Effect	Notes
cloadi	0x00	sD, cIMM	$sD \leftarrow cIMM$	cIMM packed in imm_90 (Re/Im per above).
cadd_i	0x01	sD, sA, cIMM	$sD \leftarrow sA + cIMM$	Saturating per scalar; Q32.32 truncation as needed.
cmul_i	0x02	sD, sA, cIMM	$sD \leftarrow sA \times cIMM$	Widen internally, clamp/truncate to Q32.32.
csub_i	0x03	sD, sA, cIMM	$sD \leftarrow sA - cIMM$	Saturating; truncation semantics match csub.c.
cdiv_i	0x04	sD, sA, cIMM	$sD \leftarrow sA \div cIMM$	$(sA \cdot cIMM) / cIMM ^2$; if $ cIMM =0$ then $sD:=0$.
cmaxabs_i	0x05	sD, sA, cIMM	$sD \leftarrow \arg \max_{x \in \{sA, cIMM\}} x $	Ties pick sA.
cminabs_i	0x06	sD, sA, cIMM	$sD \leftarrow \arg \min_{x \in \{sA, cIMM\}} x $	Ties pick sA.
cscale_i	0x10	sD, sA, rIMM	$sD \leftarrow sA \times rIMM$	Real scale rIMM packed in imm_90 as Q22.23 with Im=0.

J-type — Conditional Jump

OPCODE: 0x03

Table 7: J-type: Conditional jump (single predicate)

Mnemonic	subop	Operands	Effect	Notes
jrel	0x00	offs33	If $s1 \neq 0$: $PC \leftarrow PC + \text{offs33}$ (instruction-relative).	Encoding: offs33 is <i>signed</i> two's complement in <i>instruction units</i> (128-bit words); base is the address of <i>this</i> instruction. Field rs1 must be 001b (predicated on s1). All flags16 and reserved bits must be zero.

S-type — Matrix-bank Vector Load/Store (stride implicit)

OPCODE: 0x04

Table 8: S-type: Matrix-bank vector load/store

Mnemonic	subop	Operands	Effect	Notes
vld	0x00	vD, mbid, rc, idx16, len16	Load into vD the sequence: if $rc=0$: $(r=\text{idx16}, c=0..L-1)$, if $rc=1$: $(r=0..L-1, c=\text{idx16})$, where $L = \text{len16}$ if nonzero, else $L = \text{VLEN}$.	Encoding: FLAGS16.rc bit 111 (MSB of FLAGS16). $rc=0$ means row, $rc=1$ means column. Fields idx16 and len16 are unsigned 16-bit . $\text{len16}=0$ selects $L = \text{VLEN}$. Row stride = 1; column stride = len16 . Elements are 1 complex (Re then Im).
vst	0x01	vS, mbid, rc, idx16, len16	Store from vS to the same address pattern as vld.	Same encoding and field conventions as vld.

Mnemonic	subop	Operands	Effect	Notes
<code>sld.xy</code>	0x02	sD, mbid, x16, y16	Load the single element at coordinates ($r=y16$, $c=x16$) from matrix-bank <code>mbid</code> into scalar register <code>sD</code> .	Coordinates are 0-based unsigned 16-bit . Mapping: $x16 \rightarrow i16[88:73]$, $y16 \rightarrow j16[72:57]$. Element size is one 128-bit complex number (then Im). Out-of-bounds coordinates trap.
<code>sst.xy</code>	0x03	sS, mbid, x16, y16	Store scalar <code>sS</code> to the element at ($r=y16$, $c=x16$) in matrix-bank <code>mbid</code> .	Same addressing and trapping rules as <code>sld.xy</code> . Mapping: $x16 \rightarrow i16[88:73]$, $y16 \rightarrow j16[72:57]$.

Global encoding rules. Unless otherwise specified, all **reserved** fields are **must-be-zero** and all currently **unused** fields are encoded as **zero**.